

## Assignment No. 2

```
In [1]: import pandas as pd
```

```
In [5]: df = pd.read_csv("emails.csv")
```

```
In [6]: df.shape
```

```
Out[6]: (5172, 3002)
```

```
In [7]: df.head()
```

```
Out[7]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructure	military	allowing	ff	dry	Prediction
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	0	0	0	0	0	0
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	0	0	0	1	0	0
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	0	0	0	0	0	0
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	0	0	0	0	0	0
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	0	0	0	1	0	0

5 rows × 3002 columns

```
In [8]: X= df.drop(['Email No.', 'Prediction'], axis =1)
y= df['Prediction']
```

```
In [9]: X.shape
```

```
Out[9]: (5172, 3000)
```

```
In [10]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5172 entries, 0 to 5171
Columns: 3000 entries, the to dry
dtypes: int64(3000)
memory usage: 118.4 MB
```

```
In [11]: X.dtypes
```


```
Out[11]: the          int64
to          int64
ect         int64
and         int64
for         int64
...
infrastructure int64
military      int64
allowing      int64
ff            int64
dry           int64
Length: 3000, dtype: object
```

```
In [12]: set(X.dtypes)
```

```
Out[12]: {dtype('int64')}
```

```
In [14]: import seaborn as sns
sns.countplot(x=y)
```

```
Out[14]: <AxesSubplot:xlabel='Prediction', ylabel='count'>
```



```
In [15]: y.value_counts()
```

```
Out[15]: 0    3672
         1    1500
         Name: Prediction, dtype: int64
```

```
In [17]: from sklearn.preprocessing import MinMaxScaler
         scaler = MinMaxScaler()
         x_scaled = scaler.fit_transform(X)
```

```
In [18]: x_scaled
```

```
Out[18]: array([[0.        , 0.        , 0.        , ..., 0.        , 0.        ,
                0.        ],
               [0.03809524, 0.09848485, 0.06705539, ..., 0.        , 0.00877193,
                0.        ],
               [0.        , 0.        , 0.        , ..., 0.        , 0.        ,
                0.        ],
               ...,
               [0.        , 0.        , 0.        , ..., 0.        , 0.        ,
                0.        ],
               [0.00952381, 0.0530303 , 0.        , ..., 0.        , 0.00877193,
                0.        ],
               [0.1047619 , 0.18181818, 0.01166181, ..., 0.        , 0.        ,
                0.        ]])
```

```
In [19]: from sklearn.model_selection import train_test_split
         x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, random_state=0, test_size = 0.30)
```

```
In [20]: x_scaled.shape
```

```
Out[20]: (5172, 3000)
```

```
In [21]: x_train.shape
```

```
Out[21]: (3620, 3000)
```

```
In [22]: X_test.shape
```

```
Out[22]: (1552, 3000)
```

```
In [23]: from sklearn.neighbors import KNeighborsClassifier  
knn = KNeighborsClassifier(n_neighbors =5)  
knn.fit(X_train, y_train)
```

```
Out[23]: KNeighborsClassifier()
```


```
In [24]: y_pred = knn.predict(X_test)
```

C:\ProgramData\Anaconda\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: For categories with many levels, ``mode`` typically preserves the axis it a s behavior will change: the default value of ``keepdims`` will become False, the ``axis`` over which eliminated, and the value None will no longer be accepted. Set ``keepdims`` to True or False to av  
mode, \_ = stats.mode(y[neigh\_ind, k], axis=1)

```
In [29]: from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score, classification_report
```

```
In [30]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
```

```
Out[30]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x2ab10f28d60>
```



```
In [31]: y_test.value_counts()
```

```
Out[31]: 0    1111
         1     441
         Name: Prediction, dtype: int64
```

```
In [32]: accuracy_score(y_test, y_pred)
```

```
Out[32]: 0.8646907216494846
```

```
In [33]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.83	0.90	1111
1	0.69	0.95	0.80	441
accuracy			0.86	1552
macro avg	0.83	0.89	0.85	1552
weighted avg	0.90	0.86	0.87	1552

```
In [34]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [36]: error = []
for k in range(1,41):
    knn = KNeighborsClassifier(n_neighbors =k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    error.append(np.mean(y_pred != y_test))
```

```
In [37]: error
```

```
Out[37]: [0.10824742268041238,
0.10502577319587629,
0.11855670103092783,
0.11082474226804123,
0.13530927835051546,
0.12886597938144329,
0.15914948453608246,
0.15528350515463918,
0.17719072164948454,
0.17010309278350516,
0.19974226804123713,
0.19652061855670103,
0.21520618556701032,
0.21198453608247422,
0.22809278350515463,
0.22551546391752578,
0.23904639175257733,
0.23646907216494845,
0.2538659793814433,
0.25193298969072164,
0.2654639175257732,
0.26417525773195877,
0.27448453608247425,
0.27512886597938147,
0.28865979381443296,
0.2867268041237113,
0.3015463917525773,
0.3002577319587629,
0.3086340206185567,
0.30605670103092786,
0.3131443298969072,
0.3125,
0.31894329896907214,
```

```
In [38]: knn = KNeighborsClassifier(n_neighbors =1)
knn.fit(X_train, y_train)
```

```
Out[38]: KNeighborsClassifier(n_neighbors=1)
```

```
In [39]: y_pred = knn.predict(X_test)
```

```
C:\ProgramData\Anaconda\lib\site-packages\sklearn\neighbors
ions (e.g. `skew`, `kurtosis`), the default behavior of `mo
s behavior will change: the default value of `keepdims` wil
eliminated, and the value None will no longer be accepted.
    mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

```
In [40]: accuracy_score(y_test, y_pred)
```

```
Out[40]: 0.8917525773195877
```

```
In [42]: from sklearn.svm import SVC
svm = SVC(kernel = 'linear')
svm.fit(X_train, y_train)
```

```
Out[42]: SVC(kernel='linear')
```

```
In [43]: y_pred =svm.predict(X_test)
```

```
In [44]: accuracy_score(y_test, y_pred)
```

```
Out[44]: 0.9755154639175257
```